Module 7

Understanding Sqoop

Thanachart Numnonda, Executive Director, IMC Institute

Thanisa Numnonda, Faculty of Information Technology,

King Mongkut's Institute of Technology Ladkrabang

Introduction

Sqoop ("SQL-to-Hadoop") is a straightforward command-line tool with the following capabilities:

- **Imports individual tables or entire databases to files in HDFS**
- **Generates Java classes to allow you to interact with your imported data**
- **Provides the ability to import from SQL databases straight into your Hive data warehouse**

Text

See also: http://sqoop.apache.org/docs/1.4.2/SqoopUserGuide.html

Architecture Overview

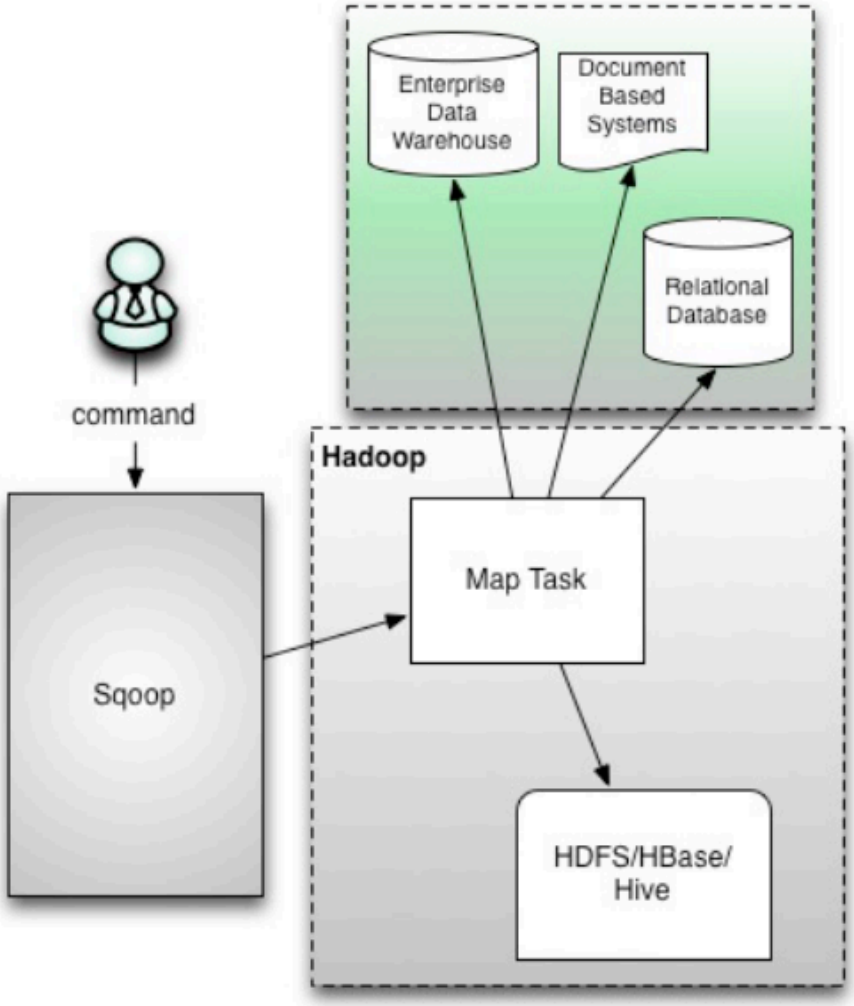

เป็น dataware housก้ได้ ไม่จำเป็นต้อง relation ขอแค่มี .....

Sqoop Benefit

- Leverages RDBMS metadata to get the column data types
- It is simple to script and uses SQL
- It can be used to handle change data capture by importing อาจจะใช้ sqoop import ไว้เห็บใน hadoop daily transactional data to Hadoop
- It uses MapReduce for export and import that enables parallel and efficient data movement

Sqoop Mode

- Sqoop import: Data moves from RDBMS to Hadoop
- Sqoop export: Data moves from Hadoop to RDBMS

Use Case #1: ETL for Data Warehouse

extract tranforms

- Transform operational data for data warehouse reports in Hadoop as the batch transformation “engine”

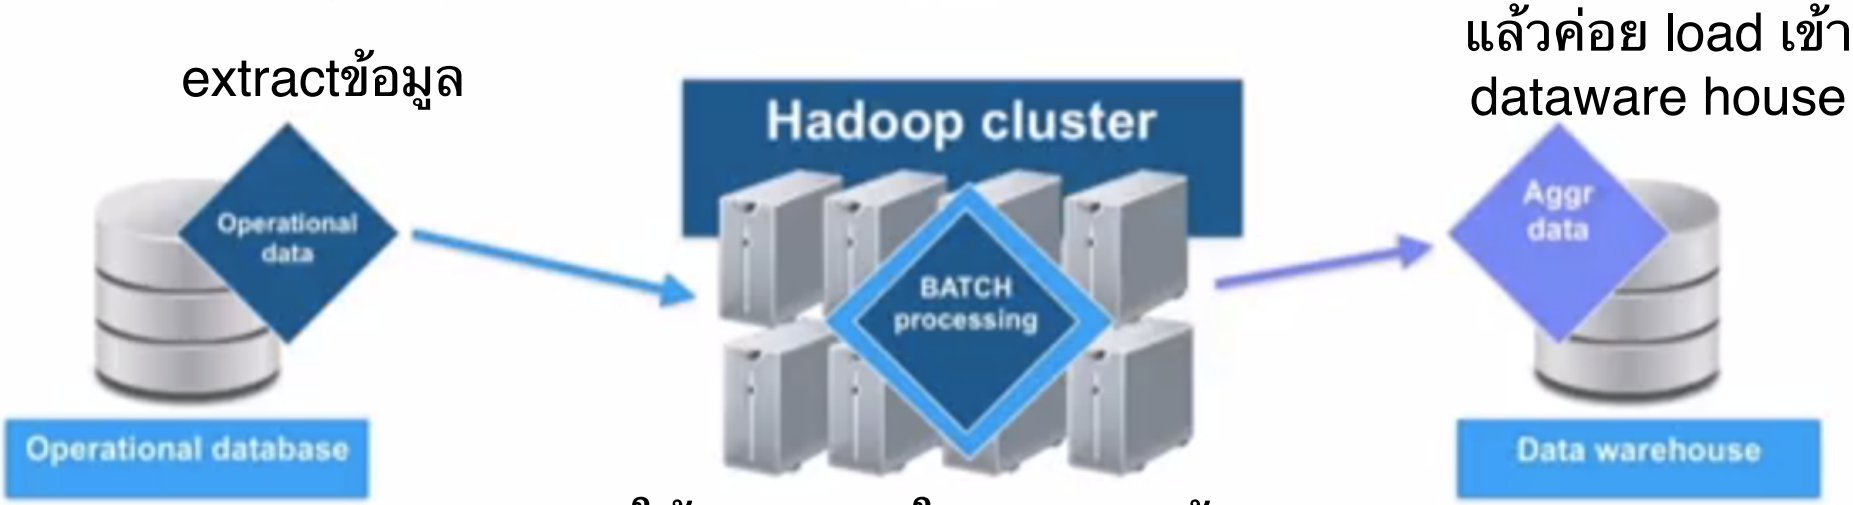

ใช้ hadoop ในการแปลงข้อมูล

Source: Mastering Apache Sqoop, David Yahalom, 2016

Use Case #2: ELT

- Extract operational data from RDBMS, process in Hadoop, return ***result*** to RDBMS

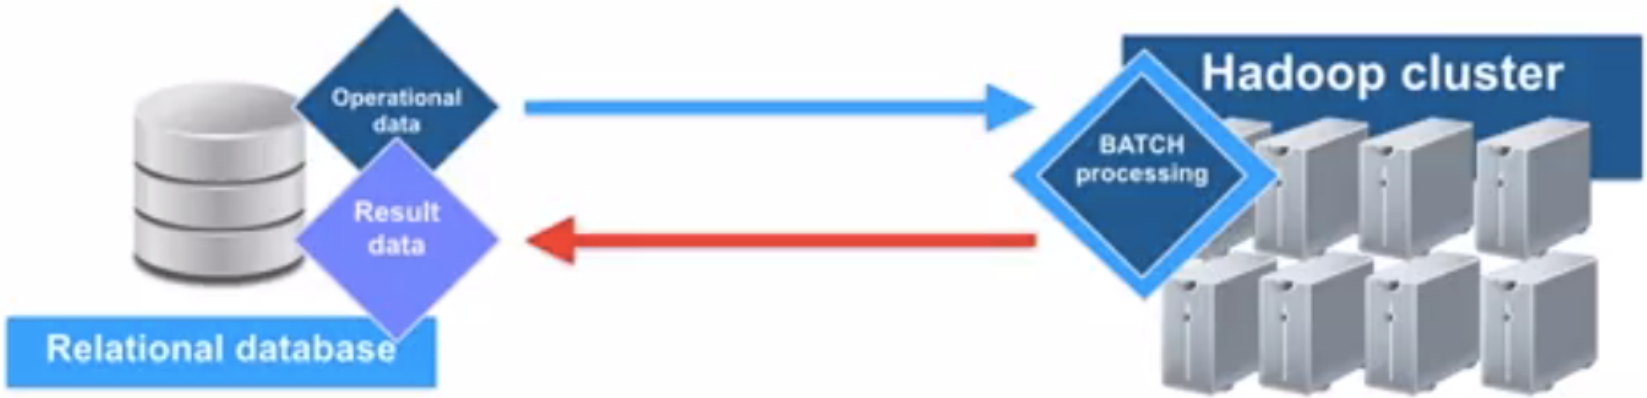

extact เข้ามาโหลดเข้ามาเก็บก่อน ค่อย transform (ส่วนมากทำแบบนี้)

Use Case #3: Data Analysis

- Copy real-time data from RDBMS, combine with raw data on Hadoop using complex data analysis logic (not just SQL!)

เพิ่มความสามารถ ในการวิเคราะห์ อาจจะมีข้อมูลดิบที่เพิ่มมาให้เราวิเคราะห์

Source: Mastering Apache Sqoop, David Yahalom, 2016

Use Case #4: Data Archival

- Move data from RDBMS after it expires to Hadoop, keeping the RDBMS “clean and lean”

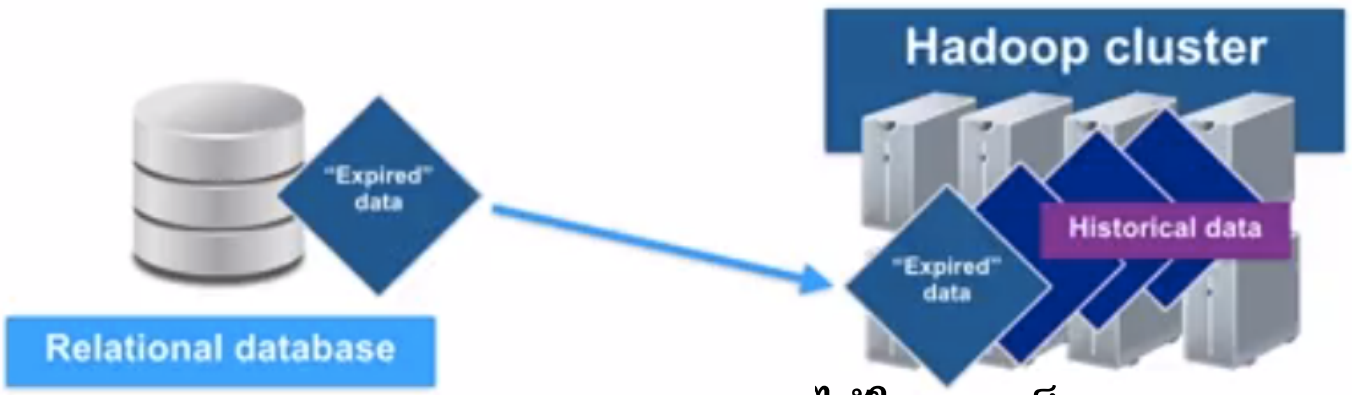

ไว้ในการเก็บ backup
เก็บเอาไว้ standby

Source: Mastering Apache Sqoop, David Yahalom, 2016

Use Case #5: Data Consolidation

- Integrate data from various organizational “data stores” to Hadoop for various data processing requirements

เอา hadoop มาเป็น consol แบบ
เก็บข้อมูลไว้หลายๆๆที่ แล้วดึงข้อมูลให้มาอยู่ที่เดียวกัน
เวลาประมาณข้อมูลจะง่ายขึ้น ใช้ sqoop ในการดึง

Source: Mastering Apache Sqoop, David Yahalom, 2016

Use Case #6: Move reports to Hadoop

- Easily allow traditional data analysis and business intelligence using Hadoop's power

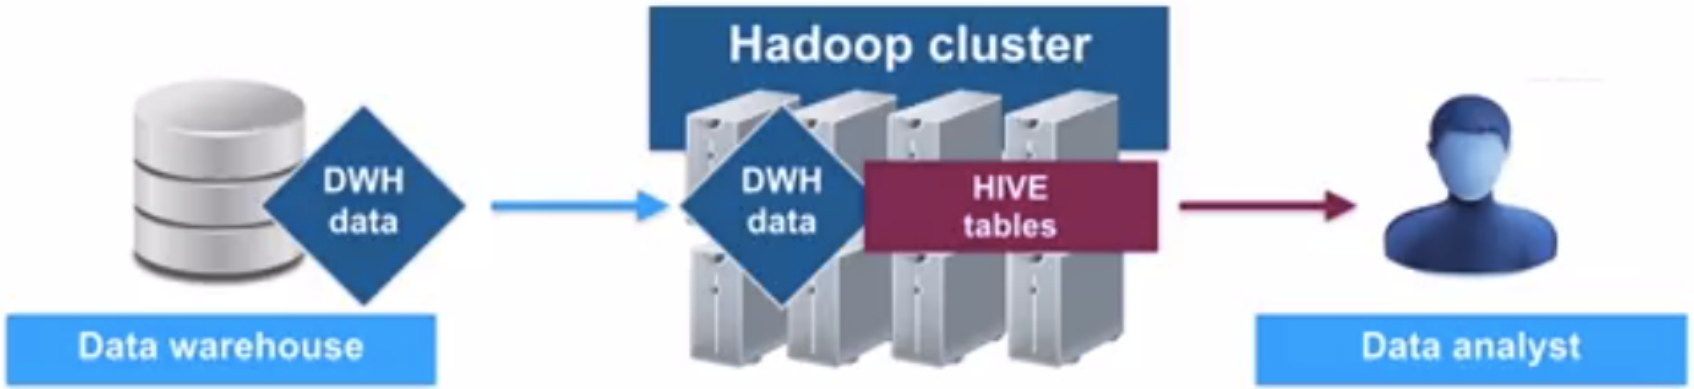

สามารถวิเคราะห์

Source: Mastering Apache Sqoop, David Yahalom, 2016

Import Commands

`--connect <jdbc-uri>`	Specifies the server or database to connect to. It also specifies the port. For example: `--connect jdbc:mysql://host:port/databaseName`
`--connection-manager <class-name>`	Specifies the connection manager class name.
`--driver <class-name>`	Specifies the fully qualified name of the JDBC driver class.
`--hadoop-home <dir>`	This parameter is used to override the `$HADOOP_HOME` environment variable.

Import Commands

`--password <password>`	Sets the authentication password required to connect to the input source.
`--username <username>`	Sets the authentication username.
`--connection-param-file <properties-file>`	Specifies the connection parameter's file.
`--help`	This option will provide the usage instructions.
`--verbose`	Prints more information during a query execution.

If a user doesn't want to specify the database password along with the command, the -P option can be used to read the password from the console.

Incremental import

Parameter/argument	Description
`--check-column <column-name>`	The value of this column is used to determine the rows to be imported during the import process.
`--incremental <incremental-type>`	Specifies the type of incremental mode. Possible values are `append` and `lastmodified`.
`--last-value <value>`	Specifies the last value or the maximum value of the `check` column from the previous import. All the records whose `check` column value is greater than the value of the `–last-value` argument will be imported to HDFS.

```
bin/sqoop import -connect jdbc:mysql://localhost:3306/db1 -username root -password password -
table student -target-dir /user/abc/student -columns "student_id,address,name"  --incremental
lastmodified -last-value "2012-11-06 19:01:35"-check-column col4
```

Export Commands

Parameters	Description
`--direct`	Use the direct mode to perform the export quickly. Note that it is only supported for MySQL.
`--export-dir<dir>`	The location of input files in HDFS.
`--table <table-name>`	Name of the output table (the RDBMS table).
`-m,--num-mappers <n>`	Refers to the number of map tasks.

Hands-On: Loading Existing Data from RDBMS to Hive

Connect to MySQL in Cloudera Quickstart

```
mysql -uroot -pcloudera
```

```
[root@quickstart /]# mysql -uroot -pcloudera
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 5069
Server version: 5.1.73 Source distribution

Copyright (c) 2000, 2013, Oracle and/or its affiliates. All
 rights reserved.

Oracle is a registered trademark of Oracle Corporation and/
or its
affiliates. Other names may be trademarks of their respecti
ve
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the curre
nt input statement.

mysql>
```

Show all MySQL DBs in Cloudera by Default

```
show databases;
```

```
mysql> show databases;
+--------------------+
| Database           |
+--------------------+
| information_schema |
| cm                 |
| firehose           |
| hue                |
| metastore          |
| mysql              |
| nav                |
| navms              |
| oozie              |
| retail_db          |
| rman               |
| sentry             |
+--------------------+
12 rows in set (0.14 sec)

mysql>
```

Change Database and Show all Tables

```
use retail_db;

show tables;
```

```
mysql> use retail_db;
Reading table information for complet
n names
You can turn off this feature to get
 -A

Database changed
mysql> show tables;
+----------------------+
| Tables_in_retail_db  |
+----------------------+
| categories           |
| customers            |
| departments          |
| order_items          |
| orders               |
| products             |
+----------------------+
6 rows in set (0.00 sec)

mysql>
```

Importing the “orders” table to Hive

import export เป็นการทำ map reduce

map reduce ตรง -m 1 เปลี่ยนเป็น -m2,3,4

```
$sqoop import  --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username root --password cloudera --table orders --hive-import
--hive-overwrite --create-hive-table -m 1
```

Data sample for orders

View more...

	orders.order_id	orders.order_date	orders.order_customer_id	orders.order_status
1	1	2013-07-25 00:00:00.0	11599	CLOSED
2	2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT
3	3	2013-07-25 00:00:00.0	12111	COMPLETE
4	4	2013-07-25 00:00:00.0	8827	CLOSED
5	5	2013-07-25 00:00:00.0	11318	COMPLETE
6	6	2013-07-25 00:00:00.0	7130	COMPLETE
7	7	2013-07-25 00:00:00.0	4530	COMPLETE
8	8	2013-07-25 00:00:00.0	2911	PROCESSING
9	9	2013-07-25 00:00:00.0	5657	PENDING_PAYMENT
10	10	2013-07-25 00:00:00.0	5648	PENDING_PAYMENT
11	11	2013-07-25 00:00:00.0	918	PAYMENT_REVIEW
12	12	2013-07-25 00:00:00.0	1837	CLOSED
13	13	2013-07-25 00:00:00.0	9149	PENDING_PAYMENT

Exercise : Use Impala to show all product names that have price = 50

```
+-----------------------------------------------+
| product_name                                  |
+-----------------------------------------------+
| Nike Adult Vapor Jet 3.0 Receiver Gloves      |
| Reebok Women's Performance Regular Motion Fit |
| Nike Men's Dri-FIT Victory Golf Polo          |
| adidas Men's Germany Home Replica White Top   |
| adidas Men's Germany Home Replica White Top   |
| Stitches Men's Pittsburgh Pirates Black Poly  |
| adidas Men's Germany Away Replica Black Top   |
| Majestic Youth Replica New York Yankees Derek |
| Majestic Youth Replica Pittsburgh Pirates And |
| Majestic Youth Replica Boston Red Sox David O |
| Majestic Youth Replica Los Angeles Dodgers Ya |
| Majestic Youth Replica New York Yankees Jacob |
+-----------------------------------------------+
Fetched 12 row(s) in 0.41s
[quickstart.cloudera:21000] >
```

Hands-On: Loading Data from RDBMS to Hadoop

Running MySQL Docker

```
root@hadoop-docker:~# docker pull mysql:5.6.40
root@hadoop-docker:~# docker run --name mysqlServer
-e MYSQL_ROOT_PASSWORD=itkmitl -p 3306:3306
-d mysql:5.6.40
root@hadoop-docker:~# docker exec -it mysqlServer bash
```

```
root@262789f9c117:/# mysql -uroot -p"itkmitl"
```

```
Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql>
```

Prepare a test database table

```
mysql> CREATE DATABASE itkmitl_db;
mysql> USE itkmitl_db;
mysql> CREATE TABLE country_tbl(id INT NOT NULL, country
VARCHAR(50), PRIMARY KEY  (id));
```

```
mysql> INSERT INTO country_tbl VALUES(1, 'USA');
mysql> INSERT INTO country_tbl VALUES(2, 'CANADA');
mysql> INSERT INTO country_tbl VALUES(3, 'Mexico');
mysql> INSERT INTO country_tbl VALUES(4, 'Brazil');
mysql> INSERT INTO country_tbl VALUES(61, 'Japan');
mysql> INSERT INTO country_tbl VALUES(65, 'Singapore');
mysql> INSERT INTO country_tbl VALUES(66, 'Thailand');
```

View data in the table

```
mysql> SELECT * FROM country_tbl;
```

```
+----+-----------+
| id | country   |
+----+-----------+
|  1 | USA       |
|  2 | CANADA    |
|  3 | Mexico    |
|  4 | Brazil    |
| 61 | Japan     |
| 65 | Singapore |
| 66 | Thailand  |
+----+-----------+
7 rows in set (0.00 sec)
```

```
mysql> exit;
```

Then exit from the container by press Ctrl-P & Ctrl-Q

Inspect imcMysql IP Address

```
root@hadoop-docker:~#  docker inspect mysqlServer | grep IPAddress
```

```
"IPAddress": "172.17.0.7",
```

Restart the Cloudera docker with linking to the MySQL Docker

```
root@hadoop-docker:~# docker run --hostname=quickstart.cloudera
--privileged=true --link mysqlServer:mysqldb -t -i -p
8888:8888 cloudera/quickstart /usr/bin/docker-quickstart
```

Case I: Importing data from MySQL to HDFS

```
[root@quickstart /]# sqoop import --connect
jdbc:mysql://172.17.0.7/itkmitl_db --username root --password
itkmitl --table country_tbl --target-dir
/user/cloudera/testtable -m 1
```

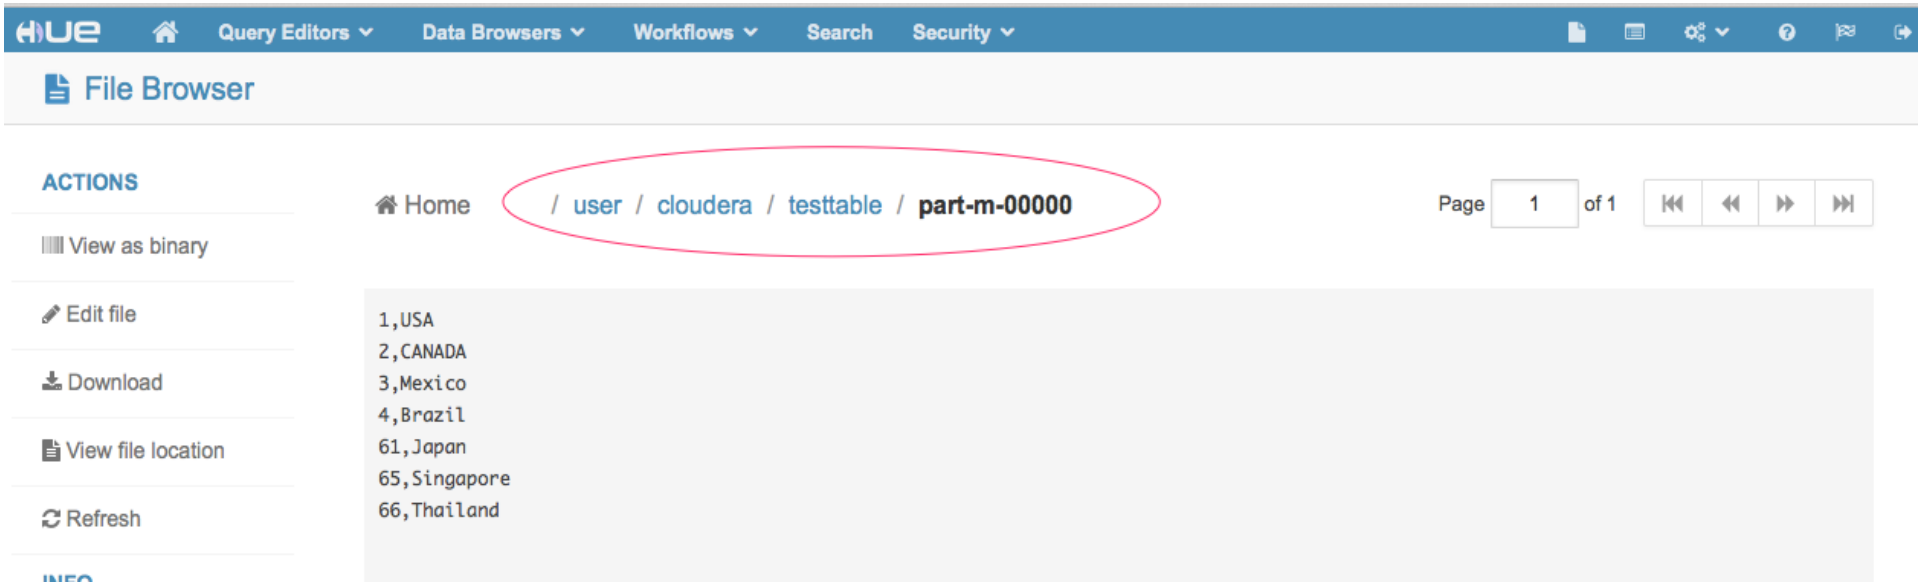

Case II: Importing data from MySQL to HBase

```
[root@quickstart /]# sqoop import --connect
jdbc:mysql://172.17.0.7/itkmitl_db --username root --password
itkmitl --table country_tbl --hbase-table country --column-
family hbase_country_cf --hbase-row-key id --hbase-create-
table -m 1
```

Start HBase

```
root@quickstart /]# hbase shell
hbase(main):001:0> list
```

```
TABLE
country
1 row(s) in 0.2570 seconds

=> ["country"]
```

Viewing HBase data

```
hbase(main):003:0> scan 'country'
ROW                     COLUMN+CELL
 1                      column=hbase_country_cf:country, timestamp=1468081466623, val
                        ue=USA
 2                      column=hbase_country_cf:country, timestamp=1468081466623, val
                        ue=CANADA
 3                      column=hbase_country_cf:country, timestamp=1468081466623, val
                        ue=Mexico
 4                      column=hbase_country_cf:country, timestamp=1468081466623, val
                        ue=Brazil
 61                     column=hbase_country_cf:country, timestamp=1468081466623, val
                        ue=Japan
 65                     column=hbase_country_cf:country, timestamp=1468081466623, val
                        ue=Singapore
 66                     column=hbase_country_cf:country, timestamp=1468081466623, val
                        ue=Thailand
7 row(s) in 0.1670 seconds
```

Viewing data from HBase browser

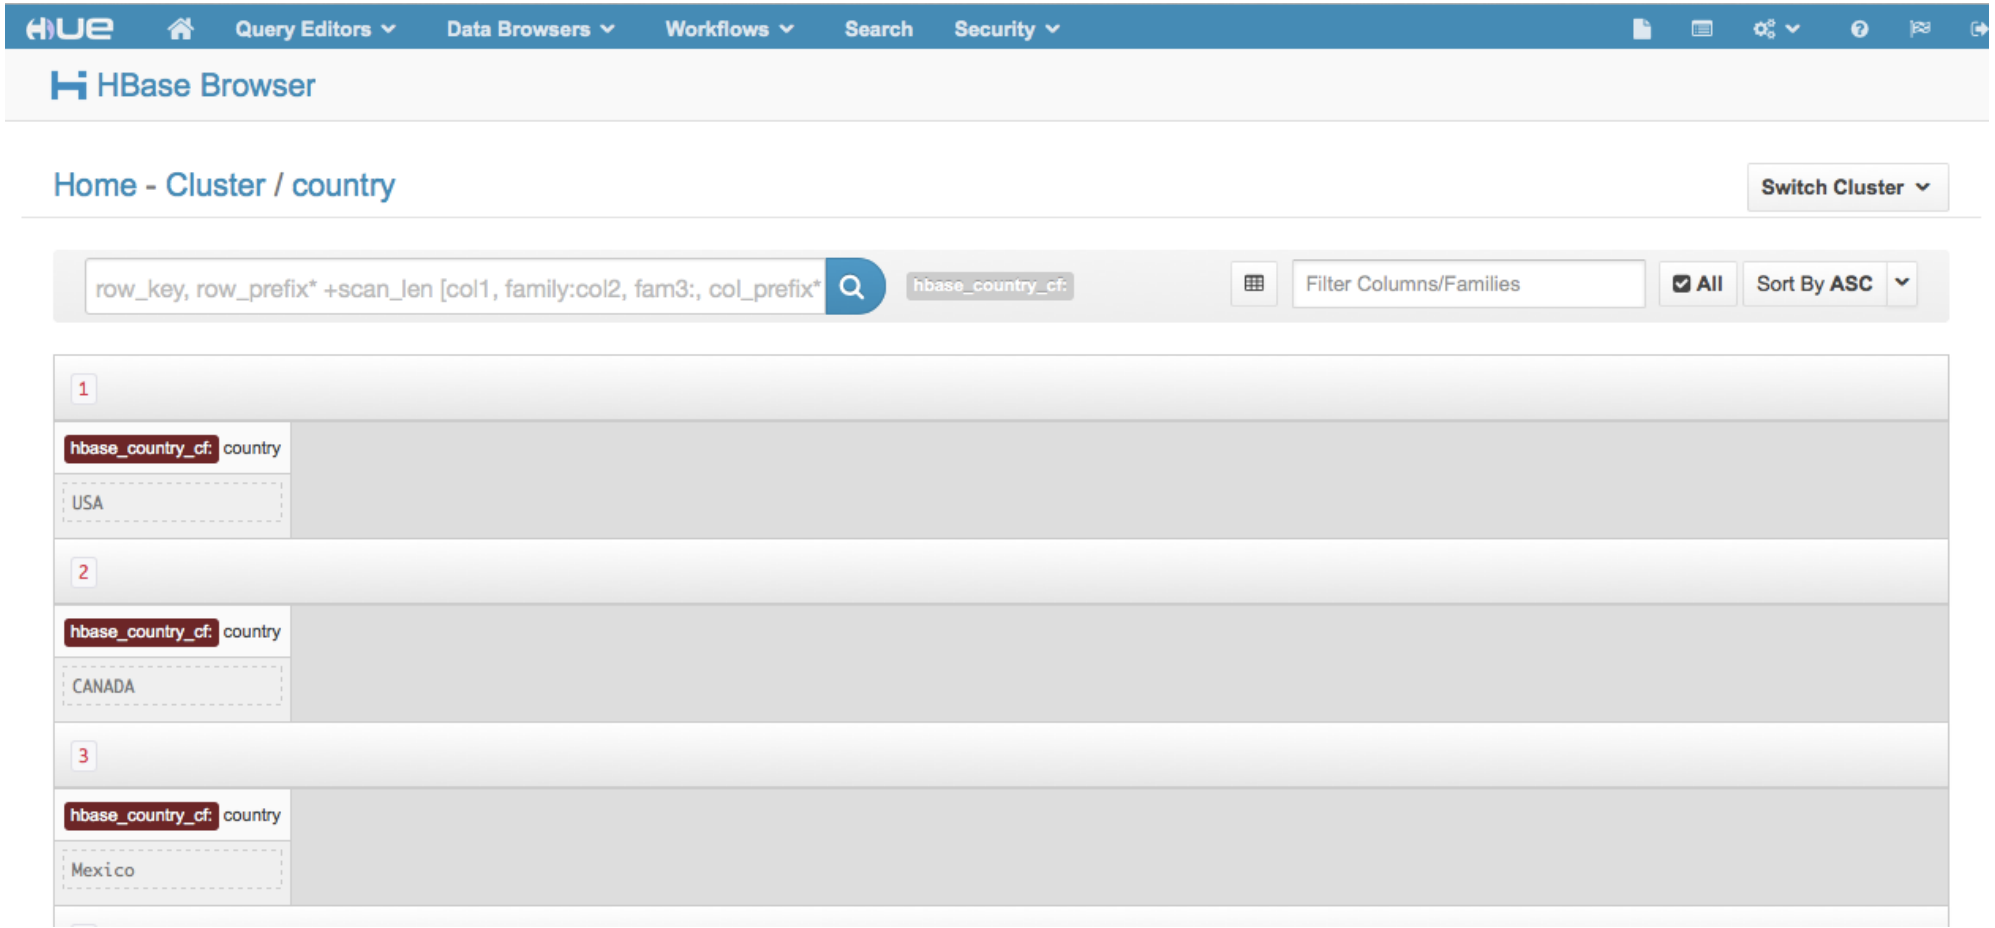

Case III: Importing data from MySQL to Hive Table

```
root@quickstart /]# sqoop import --connect
jdbc:mysql://172.17.0.7/itkmitl_db --username root --password
itkmitl --table country_tbl --hive-import --hive-table
country -m 1
```

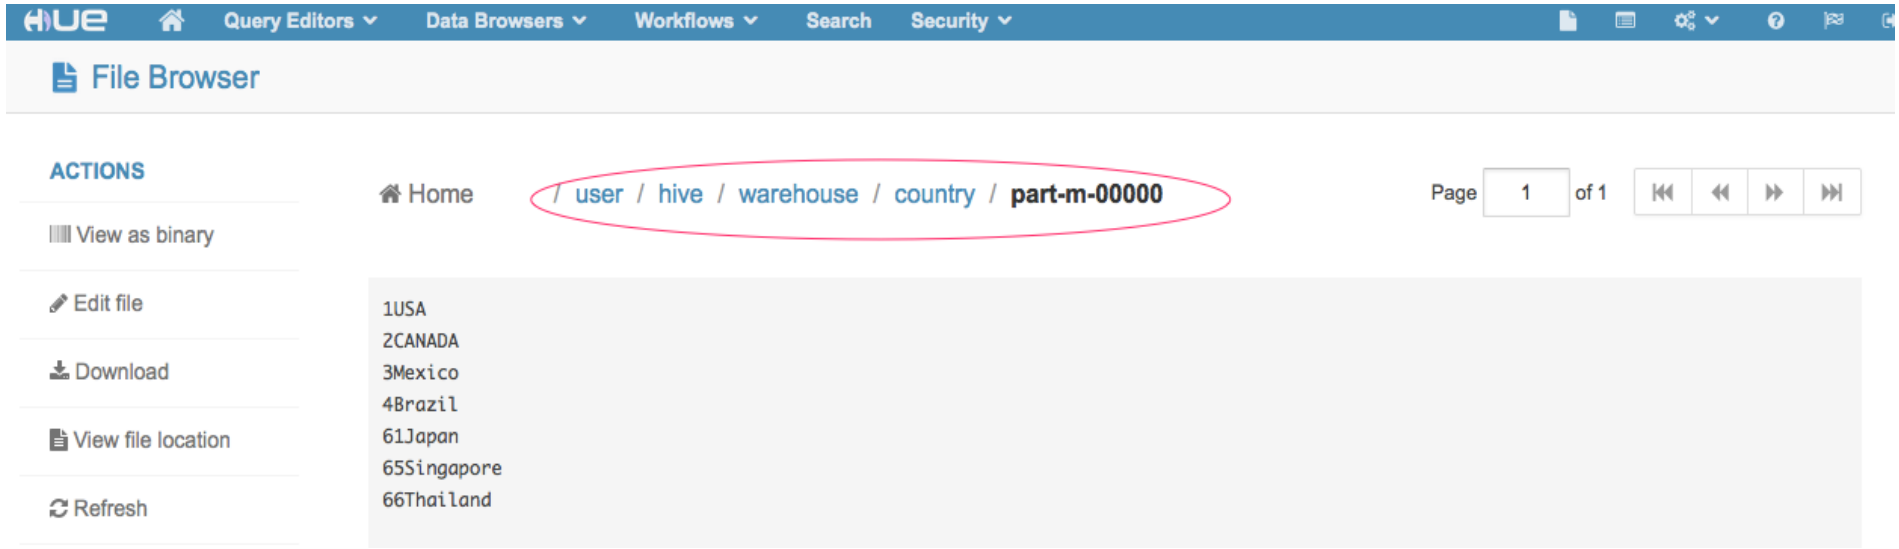

Reviewing data from Hive Table

```
[root@quickstart /]# hive
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.

hive> show tables;
hive> select * from country;
```

```
...
1       USA
2       CANADA
3       Mexico
4       Brazil
61      Japan
65      Singapore
66      Thailand
Time taken: 0.587 seconds, Fetched: 7 row(s)
```

Running from Hue: Beewax

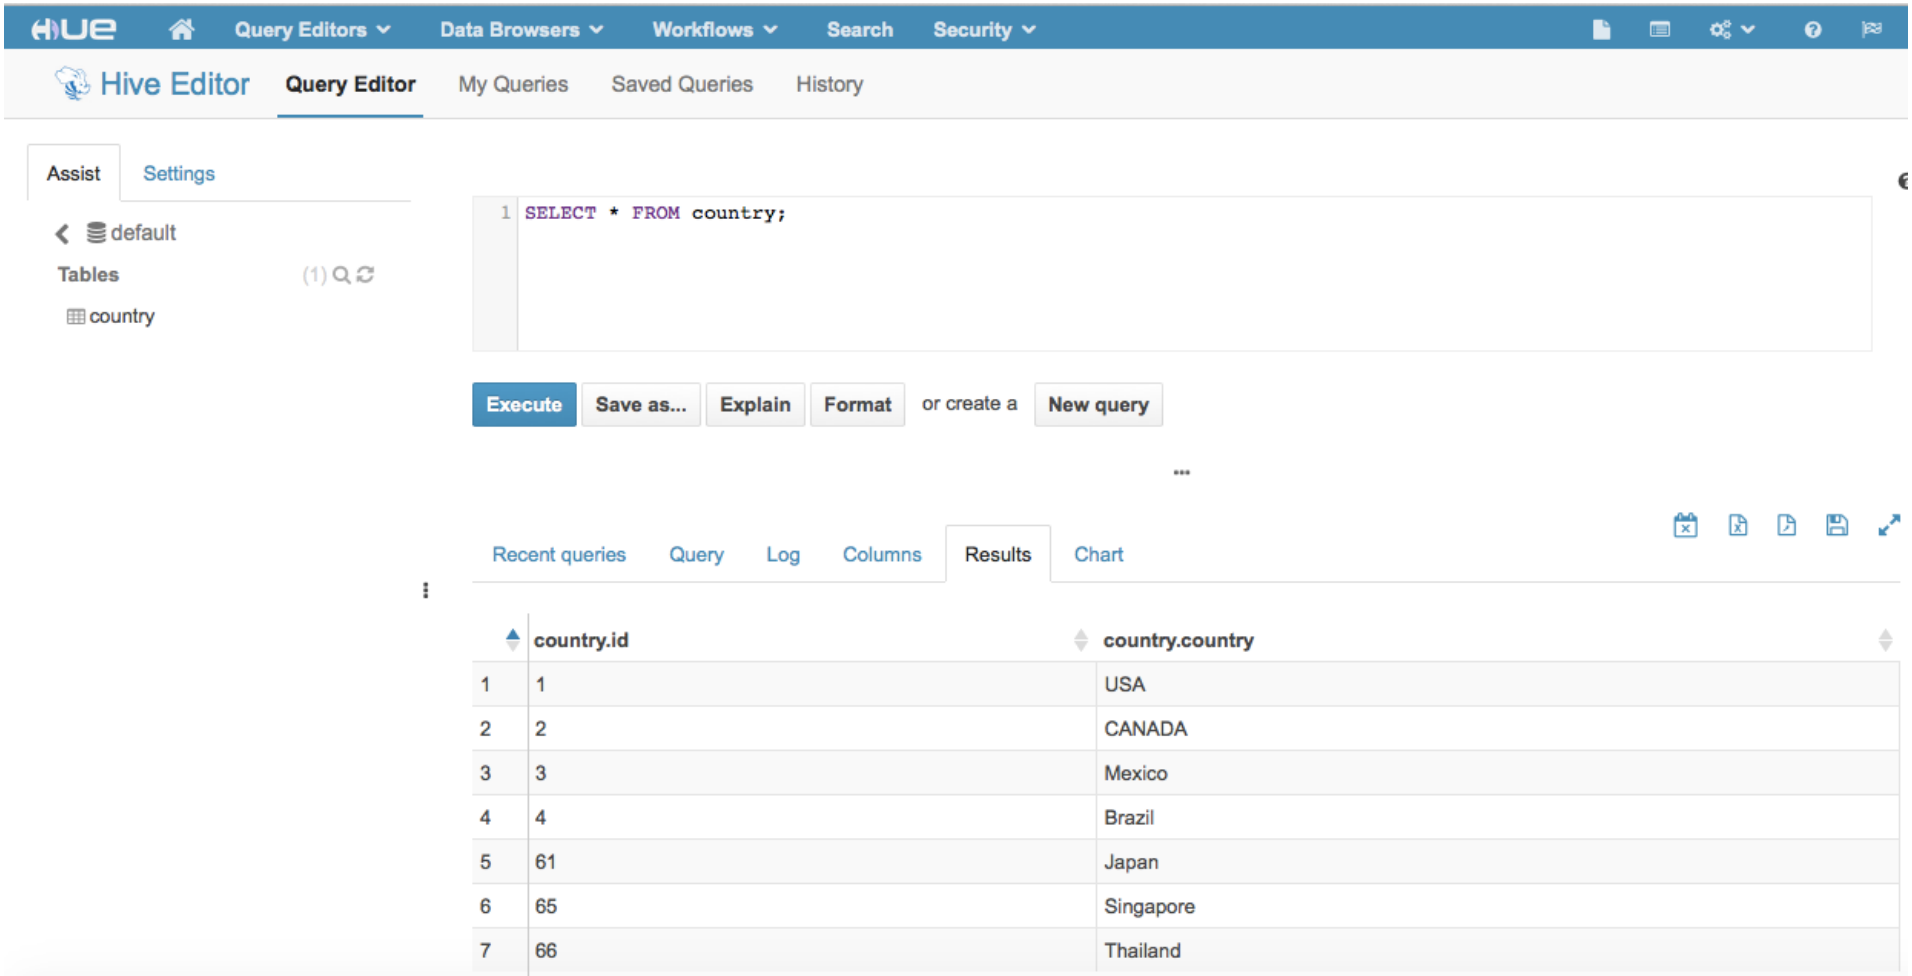

	country.id	country.country
1	1	USA
2	2	CANADA
3	3	Mexico
4	4	Brazil
5	61	Japan
6	65	Singapore
7	66	Thailand

Hands-On: Exporting Data from Hadoop to RDBMS

Exporting data from HDFS to MySQL

```
[root@quickstart /]# wget
https://s3.amazonaws.com/imcbucket/data/country.txt

[root@quickstart /]# hdfs dfs -put country.txt
/user/cloudera/input/

[root@quickstart /]# sqoop export --connect
jdbc:mysql://172.17.0.7/itkmitl_db --username root --password
itkmitl --table country_tbl --export-dir
/user/cloudera/input/country.txt
```

Viewing Result in a MySQL table on a Ubuntu server

```
root@hadoop-docker:~# docker exec -it mysqlServer bash
root@f1922a70e09c:/# mysql -uroot -p"itkmitl"
mysql> USE itkmitl_db;
mysql> SELECT * FROM country_tbl;
```